

Deep Learning: A first look

Exebit 2018, IITM

Sachin Sridhar, Department of CSE

April 15, 2018

Indian Institute of Technology Madras

Table of contents

1. Machine Learning: Some basics
2. Deep Learning
3. Feed forward networks
4. Tensorflow

Machine Learning: Some basics

What is Machine Learning?

- *“The science of getting computers to act without being explicitly programmed.”* - Andrew Ng

What is Machine Learning?

- *“The science of getting computers to act without being explicitly programmed.”* - Andrew Ng
- Uses **statistical procedures** to get computer systems to be able to “learn” and “predict” an output given an input

What is Machine Learning?

- *“The science of getting computers to act without being explicitly programmed.”* - Andrew Ng
- Uses **statistical procedures** to get computer systems to be able to “learn” and “predict” an output given an input
- Machine learning algorithms are different from conventional algorithms; they are statistical in nature and often have **no correctness guarantees**

What is Machine Learning?

- *“The science of getting computers to act without being explicitly programmed.”* - Andrew Ng
- Uses **statistical procedures** to get computer systems to be able to “learn” and “predict” an output given an input
- Machine learning algorithms are different from conventional algorithms; they are statistical in nature and often have **no correctness guarantees**
- Seen in daily life!



How does it work?

- Predicting an output given an input is mathematically formulated as a probability query $P(Y = y|X = x)$.

How does it work?

- Predicting an output given an input is mathematically formulated as a probability query $P(Y = y|X = x)$.
- We need a measure of “how good” our estimate is. So we try to minimize a *loss function* (mean squared error, negative log-likelihood, etc)

How does it work?

- Predicting an output given an input is mathematically formulated as a probability query $P(Y = y|X = x)$.
- We need a measure of “how good” our estimate is. So we try to minimize a *loss function* (mean squared error, negative log-likelihood, etc)
- This can be optimized using techniques from calculus, convex analysis, etc. A popular example of such a technique is gradient descent.

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

What could the bias of the coin be?

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

What could the bias of the coin be? (*We can never say for sure*)

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

What could the bias of the coin be? (*We can never say for sure*)

What is a good guess for the bias of the coin?

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

What could the bias of the coin be? (*We can never say for sure*)

What is a good guess for the bias of the coin?

Intuitively we would say $\frac{\text{number of heads}}{\text{total number of coin tosses}} = \frac{4}{10}$

Example 1: Coin bias

Consider the following sequence of coin tosses:

H T T H H H T T T T

What could the bias of the coin be? (*We can never say for sure*)

What is a good guess for the bias of the coin?

Intuitively we would say $\frac{\text{number of heads}}{\text{total number of coin tosses}} = \frac{4}{10}$

We can use this for predicting further coin tosses!

Coin bias: The math

Let D be the given data, and θ be a random variable representing the bias of the coin.

Definition

The *likelihood function* $L(\theta)$ is given by $L(\theta) = P(D|\theta)$.

Here we have $L(\theta) = \binom{10}{4} \theta^4 (1 - \theta)^6$ (assuming i.i.d coin tosses).

Setting the derivative wrt θ to zero to maximize the likelihood, we get:

$$\frac{dL(\theta)}{d\theta} = \binom{10}{4} (\theta^4 \cdot 6(1 - \theta)^5 + 4\theta^3 \cdot (1 - \theta)^6) = 0$$
$$\implies \boxed{\theta = 0.4}$$

In some sense, this is the best estimate we can get!

Example 2: Firm profits

Suppose we want to predict the profit of a firm given some simple data:

Units produced (x)	Cost of production (y)
10	592
20	1090
30	1604
40	2122
50	2620
70	?

We can assume that there is some fixed cost for setting up production (a), plus a marginal cost (the cost per unit) required to produce the items (b):

$$y = f(x) = a + bx$$

(This is a linear model)

Example 2: Firm profits

We require a **good** estimate for a and b .

Example 2: Firm profits

We require a **good** estimate for a and b .

Definition

Consider a dataset with n input and output pairs, denoted (x_i, y_i) . The *mean squared error* (MSE) is defined as $\frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$.

Example 2: Firm profits

We require a **good** estimate for a and b .

Definition

Consider a dataset with n input and output pairs, denoted (x_i, y_i) . The *mean squared error* (MSE) is defined as $\frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$.

In this case, the (x_i, y_i) pairs are given by $\{(10, 592), (20, 1090), (30, 1604), (40, 2122), (50, 2620)\}$. We look for values of a and b that minimize the mean squared error $L = \frac{1}{5} \sum_{i=1}^5 ((a + bx_i) - y_i)^2$.

Example 2: Firm profits

We require a **good** estimate for a and b .

Definition

Consider a dataset with n input and output pairs, denoted (x_i, y_i) . The *mean squared error* (MSE) is defined as $\frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$.

In this case, the (x_i, y_i) pairs are given by $\{(10, 592), (20, 1090), (30, 1604), (40, 2122), (50, 2620)\}$. We look for values of a and b that minimize the mean squared error $L = \frac{1}{5} \sum_{i=1}^5 ((a + bx_i) - y_i)^2$.

We can do this by setting $\frac{\partial L}{\partial a} = 0, \frac{\partial L}{\partial b} = 0$ to get $\boxed{a = 79.2, b = 50.8}$.

Firm profits: The math

$$L = \frac{1}{5} \sum_{i=1}^5 (a + bx_i - y_i)^2$$

$$\frac{\partial L}{\partial a} = \frac{1}{5} \sum_{i=1}^5 (a + bx_i - y_i) = 0 \quad (1)$$

$$\frac{\partial L}{\partial b} = \frac{1}{5} \sum_{i=1}^5 x_i (a + bx_i - y_i) = 0 \quad (2)$$

Solving, we get

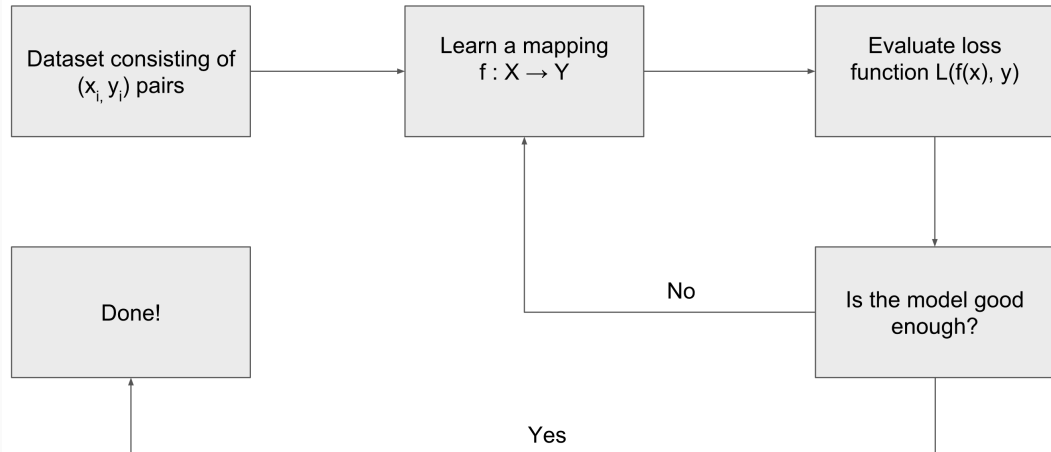
$$b = \frac{\sum_{i=1}^5 (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^5 (x_i - \bar{x})^2}, \quad a = \bar{y} - b\bar{x}$$

Where $\bar{x} = \frac{1}{5} \sum_{i=1}^5 x_i$ and $\bar{y} = \frac{1}{5} \sum_{i=1}^5 y_i$

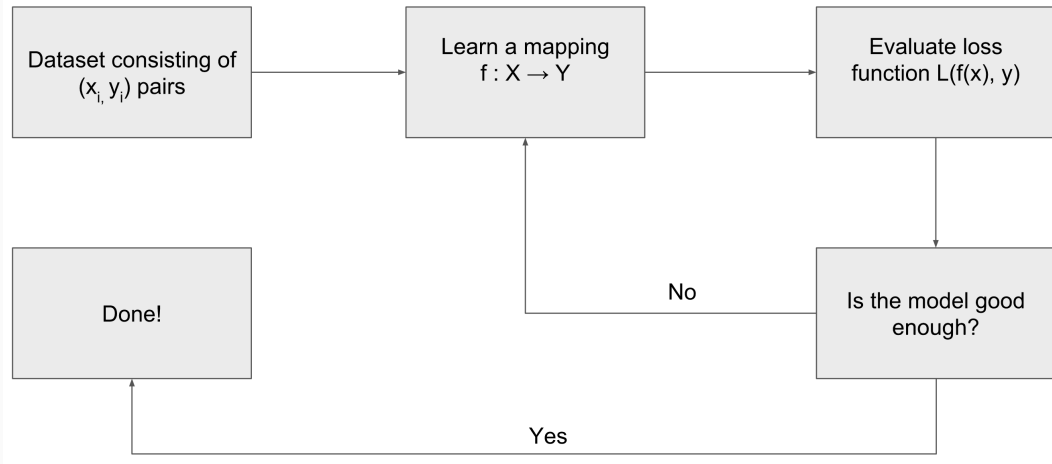
Machine Learning: A general framework

- Tries to predict an output y given an input x by learning a **mapping** $f: X \rightarrow Y$
- Quality of the estimate defined by a **loss function** $L(f(x), y)$
- ML algorithms try to minimize the loss function (closed form, gradient descent, etc)

Machine Learning: A general framework



Machine Learning: A general framework



Is it really that simple?

Machine Learning: Some issues


Overfitting: The model (function) fits too strongly to the data.

Underfitting: The model does not fit well enough to the data.

Machine Learning: Some issues

Overfitting: The model (function) fits too strongly to the data.

Underfitting: The model does not fit well enough to the data.



Deep Learning

What is Deep Learning?


- Subfield of Machine Learning inspired by a crude model of the brain, called *artificial neural networks*

What is Deep Learning?


- Subfield of Machine Learning inspired by a crude model of the brain, called *artificial neural networks*
- Typically works with large amounts of data and outperforms other machine learning algorithms

What is Deep Learning?

- Subfield of Machine Learning inspired by a crude model of the brain, called *artificial neural networks*
- Typically works with large amounts of data and outperforms other machine learning algorithms
- Can be used to learn powerful *representations* and highly non linear functions



What can Deep Learning do?



How did it all start?

- In the 1940s, Frank Rosenblatt developed an **artificial neuron model** that was capable of learning (the “perceptron”)


How did it all start?

- In the 1940s, Frank Rosenblatt developed an **artificial neuron model** that was capable of learning (the “perceptron”)
- Worked well on small problems that were *linearly separable* (what’s that?)

How did it all start?

- In the 1940s, Frank Rosenblatt developed an **artificial neuron model** that was capable of learning (the “perceptron”)
- Worked well on small problems that were **linearly separable** (what’s that?)
- Was eventually superseded by support vector machines (SVMs) which were more stable and performed better on real-world data

Perceptrons



Schematic of Rosenblatt's perceptron.

Perceptrons: The math


Mathematically, if an input $x = (x_1, \dots, x_m)$ is given to the perceptron, the output y is calculated by

$$\hat{y} = f(w_0 + w_1x_1 + \dots + w_mx_m)$$

Where $f(\cdot)$ is the *activation function*, which is a non-linear function of its input. The weights w_i are parameters of the perceptron, and they are *learned* by the perceptron.

Some commonly used activation functions are **step**, **sigmoid**, **ReLU**, **tanh**, etc.

Perceptrons: In real life



Perceptrons: What do they do?


Consider a perceptron that uses a step function for its activation:

$$\begin{aligned}\hat{y} &= f(w_0 + w_1x_1 + \dots + w_mx_m) \\ &= \begin{cases} 1, & w_0 + w_1x_1 + \dots + w_mx_m > 0 \\ 0, & \text{otherwise} \end{cases}\end{aligned}$$

What about the boundary region defined by $w_0 + w_1x_1 + \dots + w_mx_m = 0$?


It defines a *plane* in (m-1) dimensions.

Perceptrons: What do they do?



Perceptrons find a *separating plane* to try and classify the data.

Perceptrons: What do they do?



Perceptrons find a *separating plane* to try and classify the data.

Perceptrons: How do they learn?

Consider a dataset $\{(x^1, y^1), (x^2, y^2), \dots, (x^n, y^n)\}$. The perceptron learning algorithm is as follows:

Perceptrons: How do they learn?

Consider a dataset $\{(x^1, y^1), (x^2, y^2), \dots, (x^n, y^n)\}$. The perceptron learning algorithm is as follows:

1. Initialize all weights w_i to 0

Perceptrons: How do they learn?

Consider a dataset $\{(x^1, y^1), (x^2, y^2), \dots, (x^n, y^n)\}$. The perceptron learning algorithm is as follows:

1. Initialize all weights w_i to 0
2. For each training example (x^j, y^j) , calculate the output

$$\hat{y}^j = f(w_0 + w_1x_1^j + w_2x_2^j + \dots + w_mx_m^j)$$

Perceptrons: How do they learn?

Consider a dataset $\{(x^1, y^1), (x^2, y^2), \dots, (x^n, y^n)\}$. The perceptron learning algorithm is as follows:


1. Initialize all weights w_i to 0
2. For each training example (x^j, y^j) , calculate the output

$$\hat{y}^j = f(w_0 + w_1x_1^j + w_2x_2^j + \dots + w_mx_m^j)$$

3. Each of the weights w_i are updated as

$$w_i \rightarrow w_i + (y^j - \hat{y}^j)x_i^j$$

Perceptrons: Can they learn everything?



Can a linear classifier correctly classify all 4 points?

Feed forward networks

What comes next?


The perceptron is unable to learn non linear functions like XOR. Is there a way to modify it so that it can learn these functions?

What comes next?


The perceptron is unable to learn non linear functions like XOR. Is there a way to modify it so that it can learn these functions?

The idea is to **chain together** multiple perceptrons to learn more complex functions. This is done by feeding the output of a perceptron as the input of another perceptron.

Learning the XOR function




Learning the XOR function



The parameters to be learned are $w_{10}, w_{11}, w_{12}, w_{20}, w_{21}, w_{22}, w_{30}, w_{31}, w_{32}$. This successfully manages to learn the XOR function!

General feed forward networks

This idea can be extended to many layers and many perceptron units (**neurons**) in each layer.



This is called a **feed forward** neural network.

General feed forward networks

Perceptrons can represent linear functions, and a few perceptrons chained together can represent non linear functions like XOR. What could a general feed forward neural network represent?

General feed forward networks

Perceptrons can represent linear functions, and a few perceptrons chained together can represent non linear functions like XOR. What could a general feed forward neural network represent?

Theorem (Universal Approximation)

Any continuous function can be approximated to arbitrary accuracy by a feed forward neural network with the appropriate architecture.

This is great, we can learn anything we need!

General feed forward networks

Perceptrons can represent linear functions, and a few perceptrons chained together can represent non linear functions like XOR. What could a general feed forward neural network represent?

Theorem (Universal Approximation)

Any continuous function can be approximated to arbitrary accuracy by a feed forward neural network with the appropriate architecture.

This is great, we can learn anything we need! (right?)

- (Stochastic) gradient descent is used to optimize the loss function for learning.

Learning in feed forward networks

- (Stochastic) gradient descent is used to optimize the loss function for learning.
- The gradient is calculated using the *backpropagation algorithm*, a dynamic programming algorithm that keeps track of various derivatives that are required to construct the final gradient.

Learning in feed forward networks

- (Stochastic) gradient descent is used to optimize the loss function for learning.
- The gradient is calculated using the *backpropagation algorithm*, a dynamic programming algorithm that keeps track of various derivatives that are required to construct the final gradient.
- **Need not write code for this**, already done for you by software

The backpropagation algorithm (*)

- L : The total number of layers in the network (0 indexed)
- n_l : The number of neurons in the layer l
- x_j^l : The input to the i^{th} neuron in the l^{th} layer
- o_i^l : The output of the i^{th} neuron in the l^{th} layer
- W_{ij}^l : The weight connecting neuron i of layer $l - 1$ to neuron j of layer l
- b_i^l : The bias of the neuron i in layer l
- f : The activation function used at each hidden neuron (including the output layer)
- E : The error (loss) function. Here we assume that the squared error is used.

The backpropagation algorithm (*)

Since the input x_j^l only affects neurons after it and x_j^l can be written as a function of W_{ij}^l and other variables, we can write:

$$\frac{\partial E}{\partial W_{ij}^l} = \frac{\partial E}{\partial x_j^l} \frac{\partial x_j^l}{\partial W_{ij}^l}$$

$$\frac{\partial E}{\partial b_j^l} = \frac{\partial E}{\partial x_j^l} \frac{\partial x_j^l}{\partial b_j^l}$$

Defining $\delta_j^l \equiv \frac{\partial E}{\partial x_j^l}$, we have

$$\frac{\partial E}{\partial W_{ij}^l} = \delta_j^l \frac{\partial x_j^l}{\partial W_{ij}^l} \tag{1}$$

$$\frac{\partial E}{\partial b_j^l} = \delta_j^l \frac{\partial x_j^l}{\partial b_j^l} \tag{2}$$

The backpropagation algorithm (*)

By the construction of the network,

$$x_j^l = \left[\sum_{i=1}^{n_{l-1}} w_{ij}^l o_i^{l-1} \right] + b_j^l$$

And so

$$\frac{\partial x_j^l}{\partial w_{ij}^l} = o_i^{l-1} \tag{3}$$

$$\frac{\partial x_j^l}{\partial b_j^l} = 1 \tag{4}$$

The backpropagation algorithm (*)

Now, if l is an output layer we can write

$$\delta_j^l = \frac{\partial E}{\partial x_j^l} = \frac{\partial E}{\partial o_j^l} \frac{\partial o_j^l}{\partial x_j^l}$$

Since we are using the squared loss, $E = \frac{1}{2}(\sum_{j=1}^{n_L}(t - o_j^l)^2)$ where t is the target output. This gives an easy way to calculate the first term. For the second term, note that the derivative of the output of a neuron with respect to the input is just the derivative of the activation function. Putting these two together, we get

$$\delta_j^l = (t - o_j^l)f'(x_j^l), \quad \text{If } l \text{ is an output layer} \quad (5)$$

The backpropagation algorithm (*)

If l is not an output layer, we can determine δ_j^l in terms of the δ_k^{l+1} values. Since changing the input to a neuron affects all inputs of the next layer of neurons, we can write

$$\delta_j^l = \frac{\partial E}{\partial x_j^l} = \sum_{k=0}^{n_{l+1}} \frac{\partial E}{\partial x_k^{l+1}} \frac{\partial x_k^{l+1}}{\partial o_j^l} \frac{\partial o_j^l}{\partial x_j^l}$$

Finally we have $x_k^{l+1} = \sum_{j=1}^{n_l} W_{jk}^{l+1} o_j^l$, so $\frac{\partial x_k^{l+1}}{\partial o_j^l} = W_{jk}^{l+1}$, and

$$\delta_j^l = \sum_{k=0}^{n_{l+1}} \delta_k^{l+1} W_{jk}^{l+1} f'(x_j^l), \quad \text{If } l \text{ is not an output layer} \quad (6)$$

The backpropagation algorithm (*)

Putting together (1), (2), (3), (4), (5), (6), we get

$$\frac{\partial E}{\partial W_{ij}^l} = \delta_j^l o_i^{l-1}$$

$$\frac{\partial E}{\partial b_j^l} = \delta_j^l$$

$$\delta_j^l = (t - o_j^l) f'(x_j^l), \quad \text{If } l \text{ is an output layer}$$

$$\delta_j^l = \sum_{k=0}^{n_{l+1}} \delta_k^{l+1} W_{jk}^{l+1} f'(x_j^l), \quad \text{If } l \text{ is not an output layer}$$

The backpropagation algorithm (*)

This can be written in matrix/vector form as

Backpropagation equations

$$\frac{\partial E}{\partial W^l} = \delta^l (o^{l-1})^T$$

$$\frac{\partial E}{\partial b^l} = \delta^l$$

$$\delta^l = (t - o^l) \circ f'(x^l), \quad \text{If } l \text{ is an output layer}$$

$$\delta^l = (W^{l+1} \delta^{l+1}) \circ f'(x^l), \quad \text{If } l \text{ is not an output layer}$$

Update equation

$$W \leftarrow W - \eta \frac{\partial E}{\partial W}$$

The backpropagation algorithm (*)

This can be written in matrix/vector form as

Backpropagation equations

$$\frac{\partial E}{\partial W^l} = \delta^l (o^{l-1})^T$$

$$\frac{\partial E}{\partial b^l} = \delta^l$$

$$\delta^l = (t - o^l) \circ f'(x^l), \quad \text{If } l \text{ is an output layer}$$

$$\delta^l = (W^{l+1} \delta^{l+1}) \circ f'(x^l), \quad \text{If } l \text{ is not an output layer}$$

Update equation

$$W \leftarrow W - \eta \frac{\partial E}{\partial W}$$

The backpropagation algorithm (*)

Update equation

$$W \leftarrow W - \eta \frac{\partial E}{\partial W}$$

Where η is a parameter called the **learning rate**.

The backpropagation algorithm (*)

Update equation

$$W \leftarrow W - \eta \frac{\partial E}{\partial W}$$

Where η is a parameter called the **learning rate**.

This is way too complicated! There has to be an easier way!

The backpropagation algorithm (*)

Update equation

$$W \leftarrow W - \eta \frac{\partial E}{\partial W}$$

Where η is a parameter called the **learning rate**.

This is way too complicated! There has to be an easier way!

Enter **Tensorflow**.

Tensorflow

What is Tensorflow?

- An open source framework for ML/DL developed by Google
- Does all the heavy math for you, no need to code gradient descent. Specifying the **model architecture** is good enough
- Uses a **computational graph** internally to accomplish all of this. The graph can be manipulated too
- Has a Python library available (that we are going to use)

Placeholders

A **placeholder** in Tensorflow is an object that allows use to define computations on it without giving it a value. It is initialized as and when required, for example when giving it as input to a function.

Placeholders

A **placeholder** in Tensorflow is an object that allows use to define computations on it without giving it a value. It is initialized as and when required, for example when giving it as input to a function.

```
import tensorflow as tf  
x = tf.placeholder(tf.float32, (None, 1))
```

This declares a placeholder variable `x` that has size (unknown \times 1).

Variables

A **Variable** in Tensorflow is an object that holds a value, and is typically something that we are optimizing over (like the weights of a neural network). They must be initialized in code.

Variables

A **Variable** in Tensorflow is an object that holds a value, and is typically something that we are optimizing over (like the weights of a neural network). They must be initialized in code.

```
import tensorflow as tf  
m = tf.Variable(tf.zeros, ([1]))
```

This declares a Variable *m* that has size 1 (think of it like an array of size 1).

Variables

A **Variable** in Tensorflow is an object that holds a value, and is typically something that we are optimizing over (like the weights of a neural network). They must be initialized in code.

```
import tensorflow as tf  
m = tf.Variable(tf.zeros, ([1]))
```

This declares a Variable *m* that has size 1 (think of it like an array of size 1).

In short, **placeholders** are input data that does not change as the model trains, and **Variables** are model parameters that change with time.

A **Session** in Tensorflow is where code actually gets executed. Outside of a session, only variables and the computational graph is defined. Code execution requires that a Tensorflow Session be declared.

Sessions

A **Session** in Tensorflow is where code actually gets executed. Outside of a session, only variables and the computational graph is defined. Code execution requires that a Tensorflow Session be declared.

```
import tensorflow as tf
sess = tf.Session() # Open a session
sess.run(something) # Do stuff
sess.close() # Close the session
```

Or,

```
import tensorflow as tf
with tf.Session() as sess: # Open a session
    sess.run(something) # Do stuff
```


Firm costs revisited

Units produced (x)	Cost of production (y)
10	592
20	1090
30	1604
40	2122
50	2620

How would we find a good linear model for this problem using Tensorflow?

Firm costs revisited

(Download this at goo.gl/5yayn5)

```
1  import numpy as np
2  import tensorflow as tf
3
4  # Model linear regression  $y_{\text{hat}} = a + bx$ 
5
6  # Placeholder variables for the input data
7  x = tf.placeholder(tf.float32, [None, 1])
8  y = tf.placeholder(tf.float32, [None, 1])
9
10 # Model parameters are Variables in Tensorflow
11 a = tf.Variable(tf.zeros([1]))
12 b = tf.Variable(tf.zeros([1, 1]))
13
14 # Calculating the prediction of the linear model
15 y_hat = a + tf.matmul(x,b)
16
17 # Loss function =  $\sum((y_{\text{hat}}-y)^2)$ 
18 cost = tf.reduce_mean(tf.square(y_hat - y))
19
20 # Training using Gradient Descent to minimize cost
21 train_step = tf.train.GradientDescentOptimizer(0.000001).minimize(cost)
```

Firm costs revisited

```
22
23 sess = tf.Session()
24 init = tf.global_variables_initializer()
25 sess.run(init)
26 steps = 20000
27 for i in range(steps):
28     # Input data for the model
29     xs = np.array([[10], [20], [30], [40], [50]])
30     ys = np.array([[592], [1090], [1604], [2122], [2620]])
31
32     # Defining a feed dictionary; this is the set of
33     # (placeholder : value) pairs for calculations
34     feed = { x : xs, y : ys }
35
36     # Running a single training step
37     sess.run(train_step, feed_dict = feed)
38
39     # Printing the results
40     print("a: %f" % sess.run(b))
41     print("b: %f" % sess.run(a))
42
```

Computational graph

```
1  import numpy as np
2  import tensorflow as tf
3
4  # Model linear regression  $y_{\text{hat}} = a + bx$ 
5
6  # Placeholder variables for the input data
7  x = tf.placeholder(tf.float32, [None, 1])
8  y = tf.placeholder(tf.float32, [None, 1])
9
10 # Model parameters are Variables in Tensorflow
11 a = tf.Variable(tf.zeros([1]))
12 b = tf.Variable(tf.zeros([1, 1]))
13
14 # Calculating the prediction of the linear model
15 y_hat = a + tf.matmul(x, b)
16
17 # Loss function =  $\sum((y_{\text{hat}} - y)^2)$ 
18 cost = tf.reduce_mean(tf.square(y_hat - y))
19
20 # Training using Gradient Descent to minimize cost
21 train_step = tf.train.GradientDescentOptimizer(0.000001).minimize(cost)
```

A few important things about Tensorflow

- *All* computations happen within a session. This includes printing the value of a variable!

A few important things about Tensorflow

- *All* computations happen within a session. This includes printing the value of a variable!
- To build the computational graph, all arithmetic operations for the model must be done by Tensorflow. For example, `tf.reduce_sum` must be used instead of Python's builtin `sum()` function.

A few important things about Tensorflow

- *All* computations happen within a session. This includes printing the value of a variable!
- To build the computational graph, all arithmetic operations for the model must be done by Tensorflow. For example, `tf.reduce_sum` must be used instead of Python's builtin `sum()` function.
- When declaring a placeholder, if a dimension is “None”, it simply means that it is unknown.

Thank you!